

Optimizing Data Storage: A Parquet Benchmarking Study

Evaluating Parquet Runtime and Output Size Across Diverse Compression and Encoding Formats

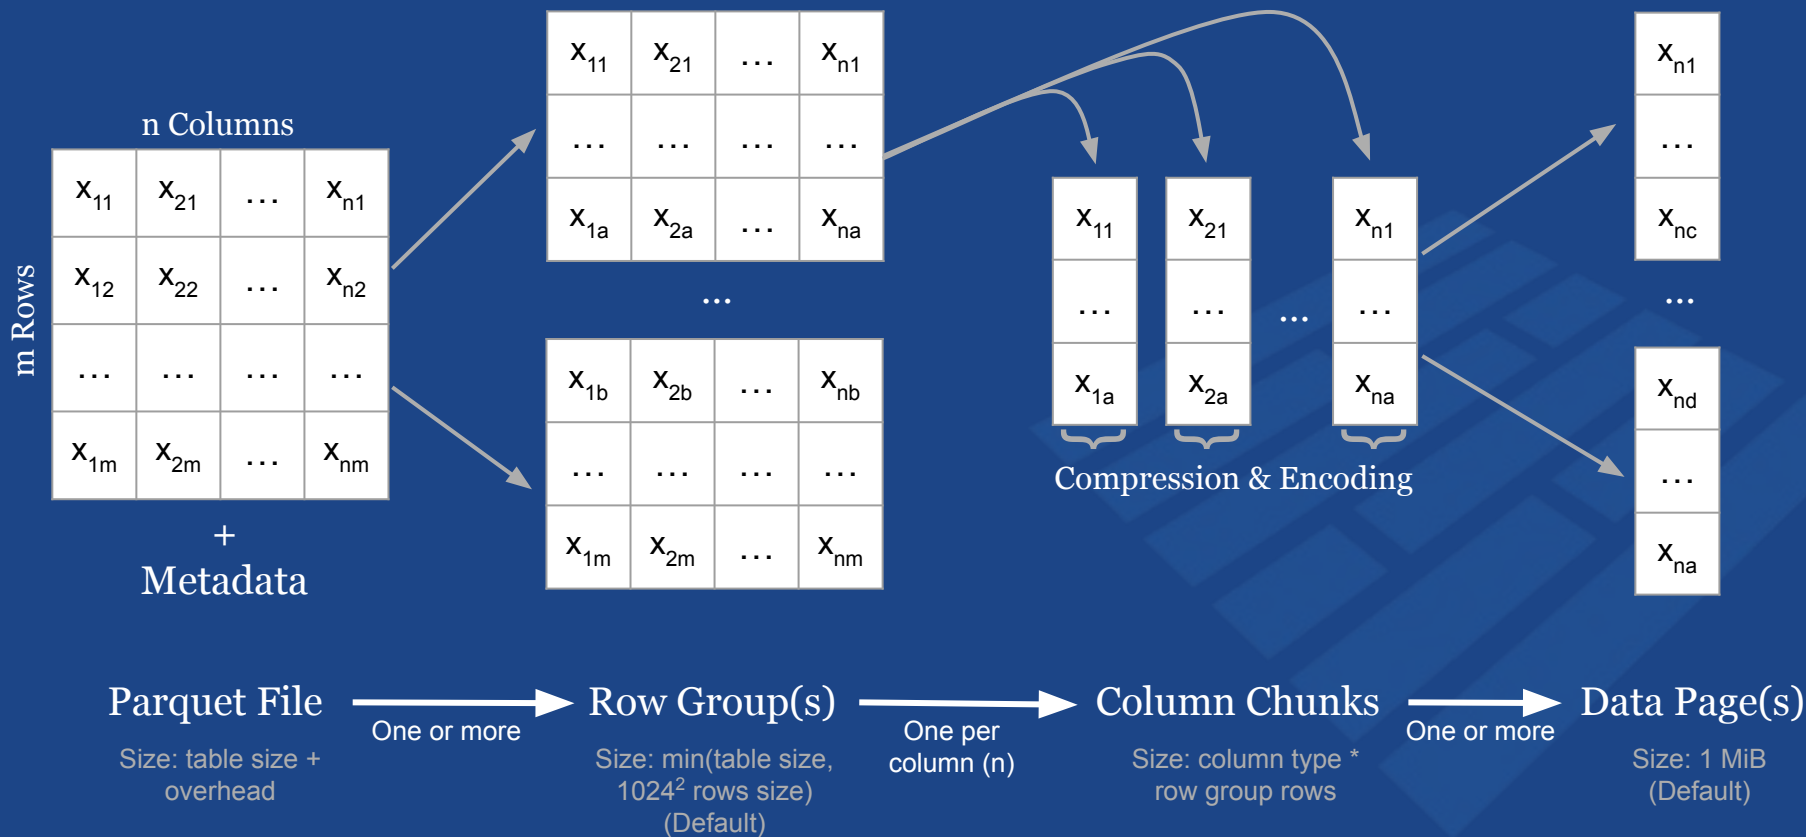
Josh Lim

August 27th, 2024 LA-UR-24-31733



Parquet Structure

Columnar Storage Hierarchy

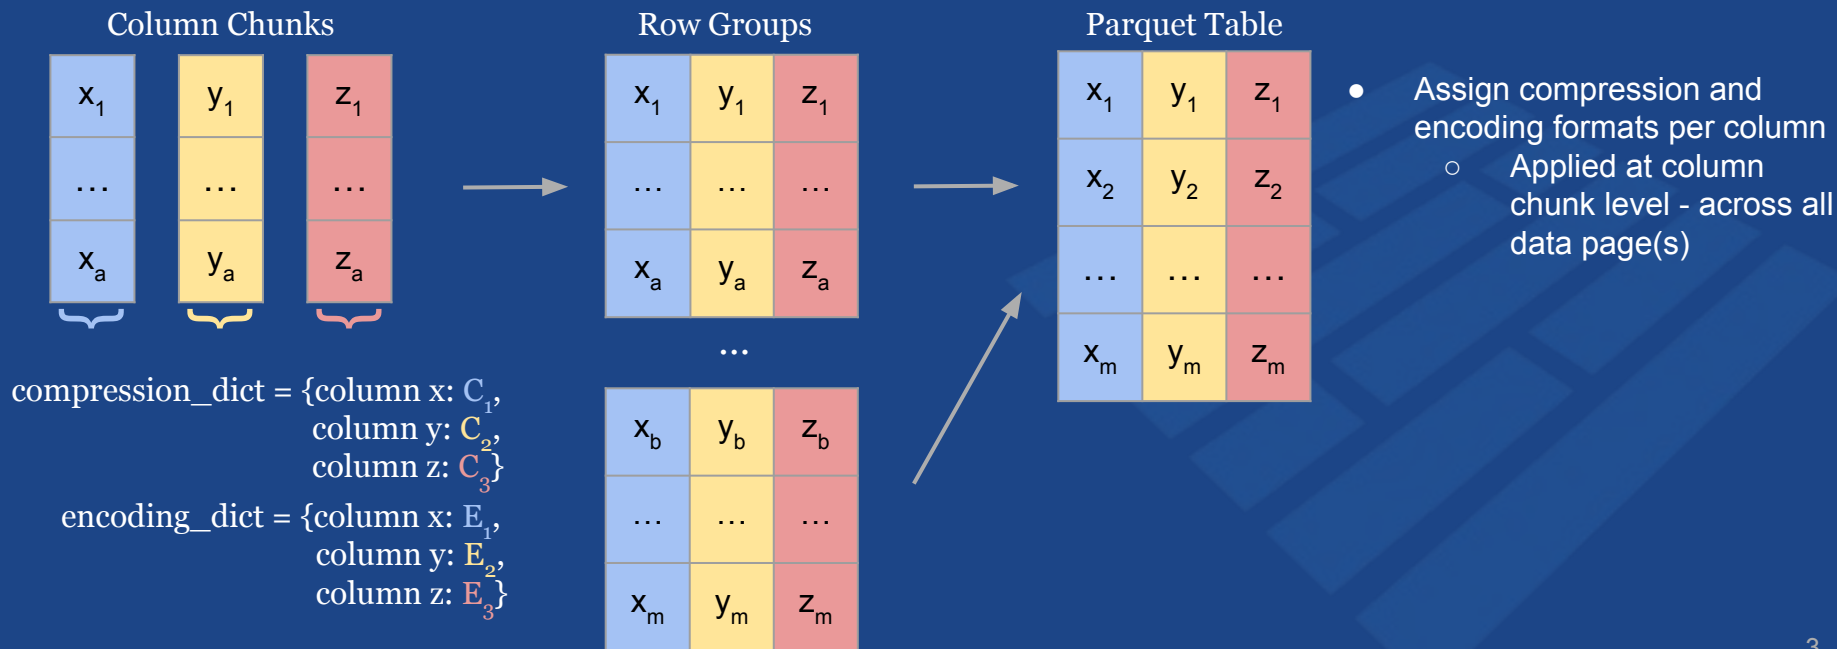


Column Based Implementation

Columnar Storage Benefits

```
parquet.write_table(data, compression=compression_dict, col_encoding=encoding_dict)
```

— Schema: x, y, z = various data types



Supported Formats

Columnar Storage Benefits

Compressions

- None
- Snappy
- Brotli
- Gzip
- ZSTD
- LZ4

Encodings

- Plain
- Byte Stream Split
- RLE/Bit-Packing
- Plain Dictionary
- RLE Dictionary
- Delta Byte Array
- Delta Binary Packed
- Delta Length Byte Array

Efficient Query and Retrieval

Columnar Storage Benefits

Parquet File

- Header (magic number “PAR1”)
- Row Group 1
 - Column Chunk: Column 1
 - ...
 - Column Chunk: Column n
- ...
- Row Group r
 - Column Chunk: Column 1
 - ...
 - Column Chunk: Column n
- Footer
 - Metadata (version, schema, row group offsets, etc.)
 - Row Group Metadata
 - Number of rows
 - Total byte size
 - Column Chunk Metadata
 - Statistics (min/max values, null count, etc.)
 - Format (codec, encoding, etc.)
 - Offset to Metadata

Projection

x_{11}		x_{n1}
...		...
x_{1m}		x_{nm}

- Access subset of columns.
- Reads select column chunks across all row groups from disk.

Selection

x_{1y}	...	x_{ny}
x_{1m}	...	x_{nm}

- Filter rows by comparing predicate to row group metadata statistics.
- Reads select row groups from disk.

VPIC Data

LANL C2-VPIC Sample Dataset

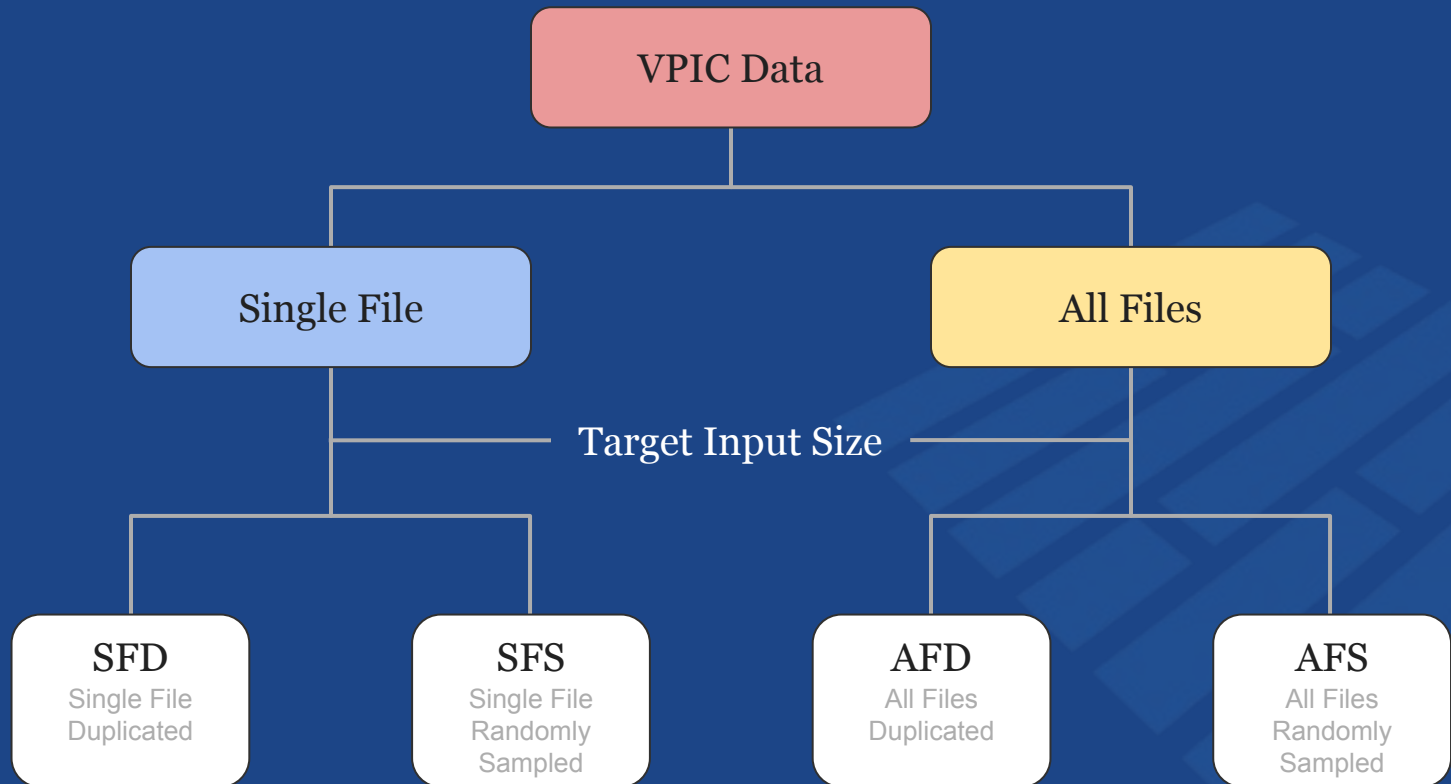
Particle Schema

8 bytes	uint64_t	ID	unique ID of a particle
8 bytes	uint64_t	padding	
4 bytes	float	x	location of particle in X direction
4 bytes	float	y	location of particle in Y direction
4 bytes	float	z	location of particle in Z direction
4 bytes	float	i	index of the cell that had the particle
4 bytes	float	ux	momentum of particle in X direction
4 bytes	float	uy	momentum of particle in Y direction
4 bytes	float	uz	momentum of particle in Z direction
4 bytes	float	ke	kinetic energy of particle

- 48 bytes per particle
- 128 * 1024 particles per file
- 6 MiB per file
- 42 total files
- 252 total MiB

VPIC Input Data Format

Benchmark Set Up



VPIC Benchmark Framework


Writing Input Data Into Parquet

SFD

SFS

AFD

AFS



- Compressions: None, Snappy, Gzip, Brotli, LZ4, ZSTD
- Encodings: Plain, Plain Dictionary, RLE Dictionary, Byte Stream Split


Vector Norm

Analyzing Benchmark Results

$$\sqrt{\left(\frac{\text{Output Size}_i}{\max(\text{Output Size})}\right)^2 + \left(\frac{\text{Average Runtime}_i}{\max(\text{Average Runtime})}\right)^2}$$


Single File Duplicated Results

Benchmark Analysis With VPIC Data




Single File Sampled Results

Benchmark Analysis With VPIC Data




All Files Duplicated Results

Benchmark Analysis With VPIC Data



All Files Sampled Results

Benchmark Analysis With VPIC Data



VPIC Findings

Parquet Benchmark Analysis

Input Format	Lowest Vector Norm	
	Compression	Encoding
SFD	ZSTD	Byte Stream Split
SFS	ZSTD	Plain Dictionary
AFD	ZSTD	Byte Stream Split
AFS	ZSTD	Byte Stream Split

Laghos Data

LANL OCS Laghos Sample Dataset

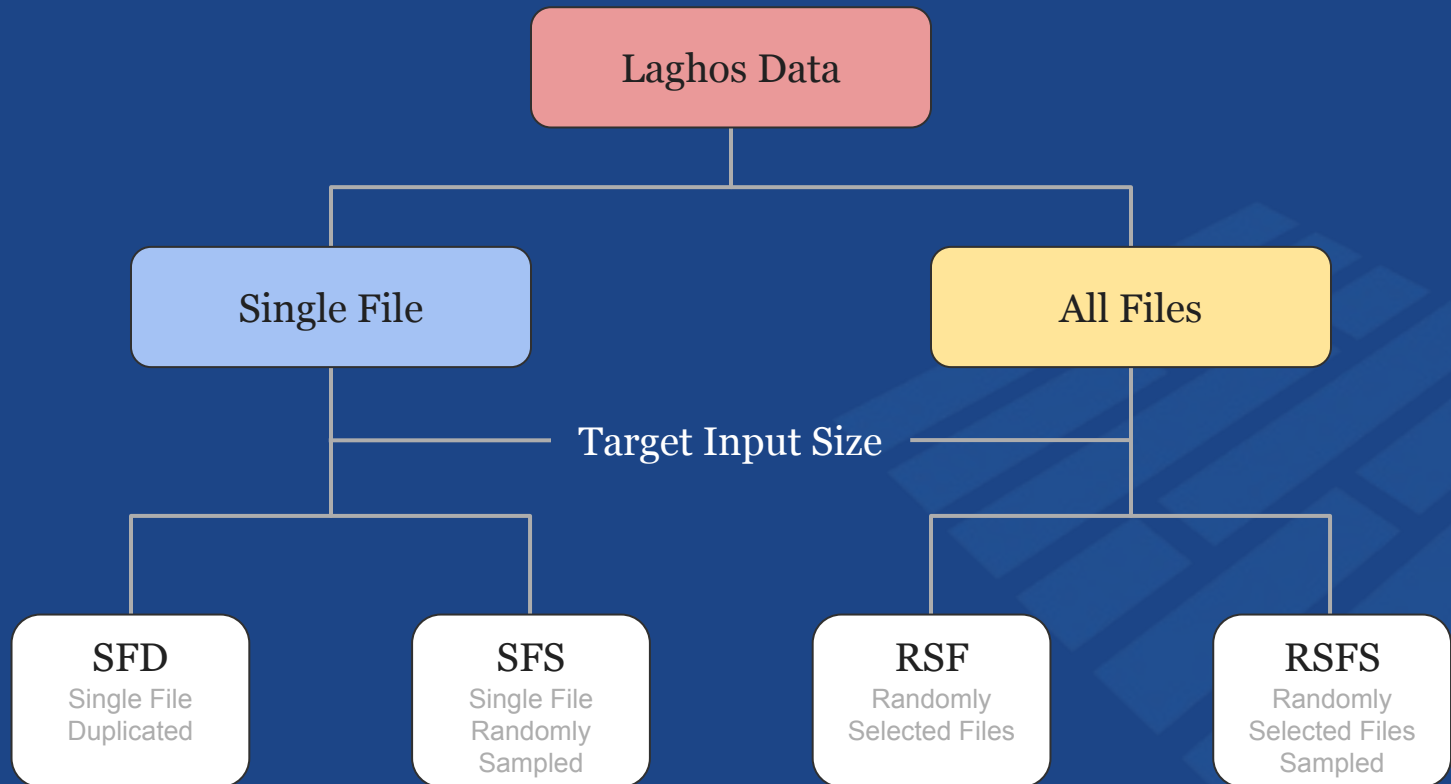
Nodal Schema

4 bytes	int32	element_id	unique ID of an element
4 bytes	int32	vertex_id	unique ID of a vertex
8 bytes	double	v_x	velocity of node in X direction
8 bytes	double	v_y	velocity of node in Y direction
8 bytes	double	v_z	velocity of node in Z direction
8 bytes	double	rho	density of node
8 bytes	double	e	energy of node
8 bytes	double	x	location of node in X direction
8 bytes	double	y	location of node in Y direction
8 bytes	double	z	location of node in Z direction

- 72 bytes per element
- 2048^2 elements per file
- 302 MiB per file
- 256 total files
- 75.5 GiB total size

Laghos Input Data Format

Benchmark Set Up



Laghos Benchmark Framework


Writing Input Data Into Parquet

SFD

SFS

RSF


RSFS



- Compressions: None, Snappy, Gzip, Brotli, LZ4, ZSTD
- Encodings: Plain, Plain Dictionary, RLE Dictionary, Byte Stream Split

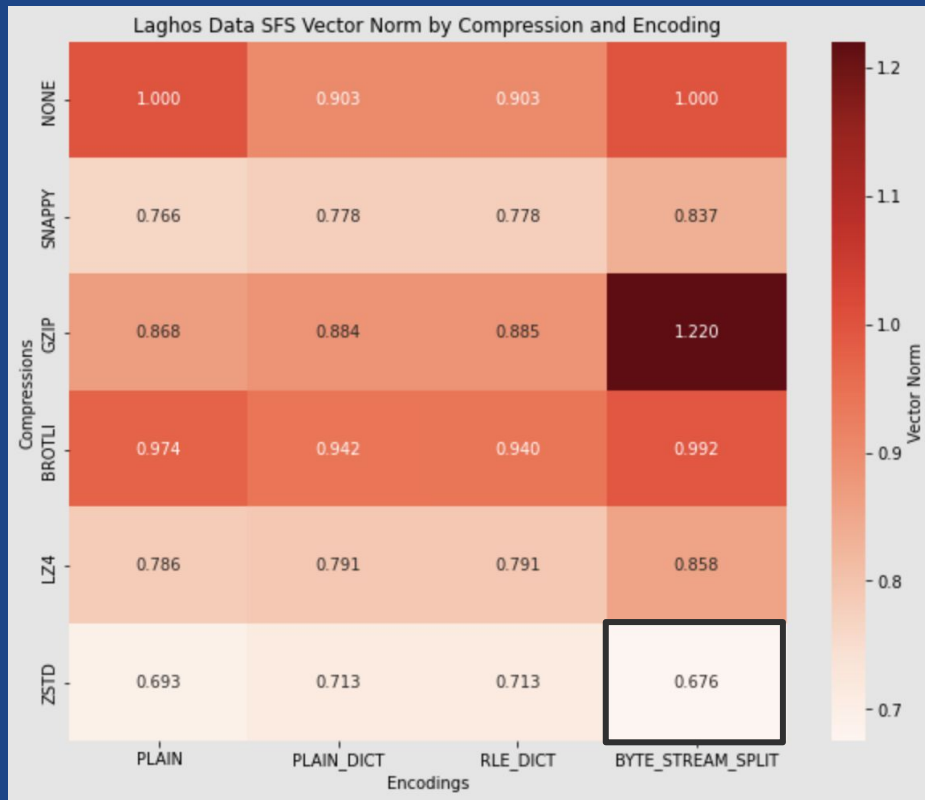
Single File Duplicated Results

Benchmark Analysis With Laghos Data




Single File Sampled Results

Benchmark Analysis With Laghos Data




Randomly Selected Files Results

Benchmark Analysis With Laghos Data



Randomly Selected Files Sampled Results

Benchmark Analysis With Laghos Data



Laghos Findings

Parquet Benchmark Analysis

Input Format	Lowest Vector Norm	
	Compression	Encoding
SFD	ZSTD	Plain
SFS	ZSTD	Byte Stream Split
RSF	ZSTD	Plain
RSFS	ZSTD	Byte Stream Split

Final Findings

Parquet Benchmark Analysis

Input Format	Lowest Vector Norm (Laghos)		Lowest Vector Norm (VPIC)	
	Compression	Encoding	Compression	Encoding
SFD	ZSTD	Plain	ZSTD	Byte Stream Split
SFS	ZSTD	Byte Stream Split	ZSTD	Plain Dictionary
RSF	ZSTD	Plain		
RSFS	ZSTD	Byte Stream Split		
AFD			ZSTD	Byte Stream Split
AFS			ZSTD	Byte Stream Split

Future Works

Parquet Benchmark Analysis

- Investigate Parquet's fallback mechanism from dictionary to plain encoding when the dictionary grows too large.
 - Analyze the performance and storage impact of switching encodings based on dictionary size or distinct value count.
 - Explore optimization opportunities by isolating conditions that trigger the fallback, aiming to enhance encoding strategies for diverse datasets.

